



Dive into PlazmaDB Performance Characteristics And Future Challenges

2018/10/16 TD tech talk

Keisuke Suzuki
Software engineer

Who am I?

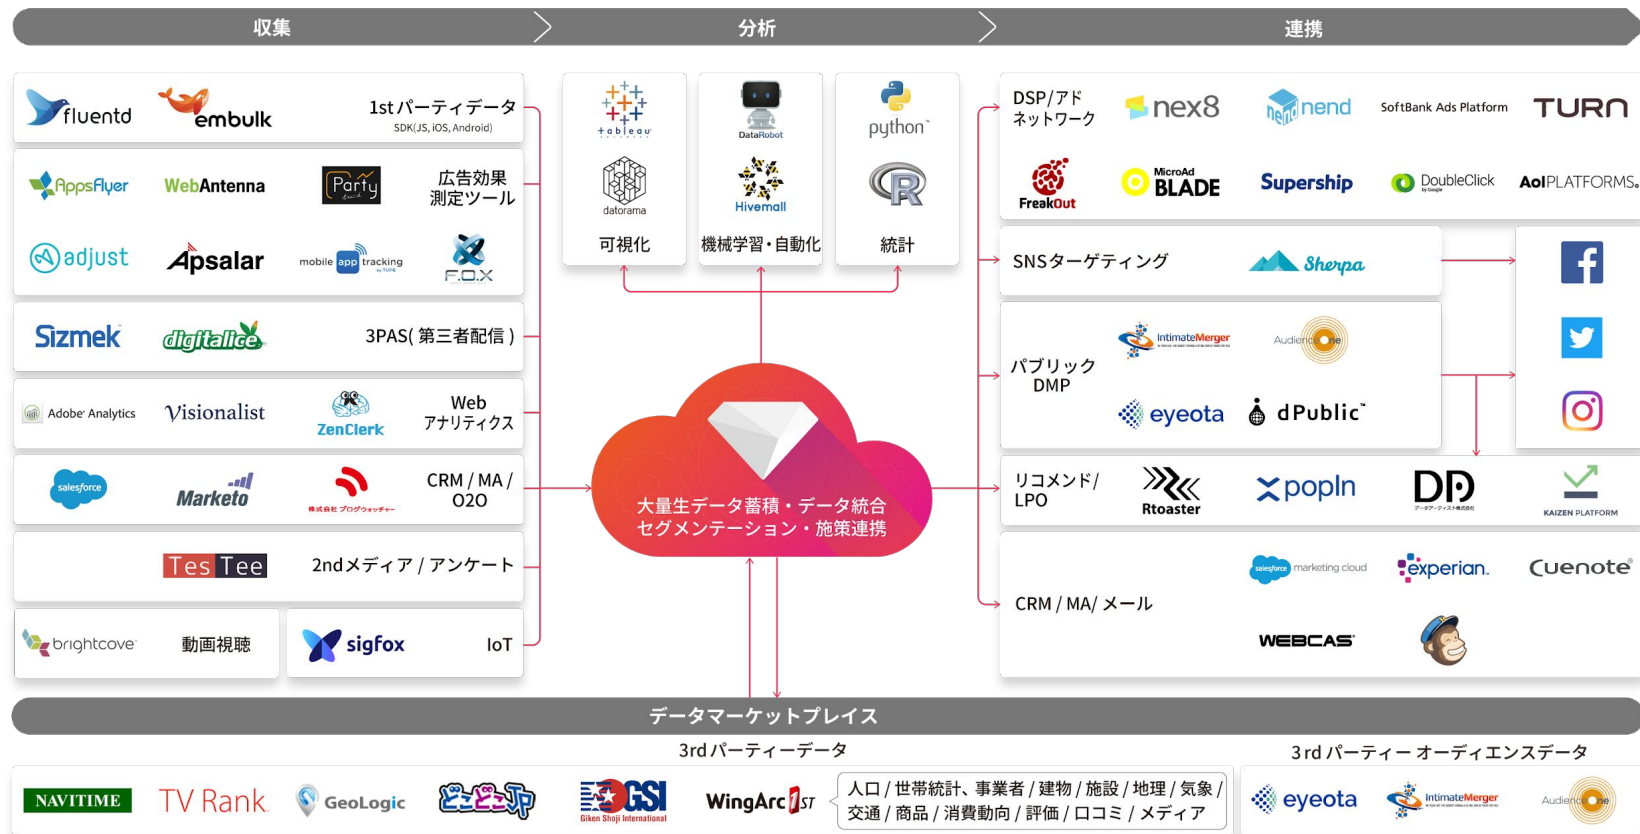
Keisuke Suzuki

- Backend Engineer @ Treasure Data KK
 - PlazmaDB: distributed storage
 - Datatank: data mart
- DB / Distributed system / Performance optimization
- Twitter: @yajilobee

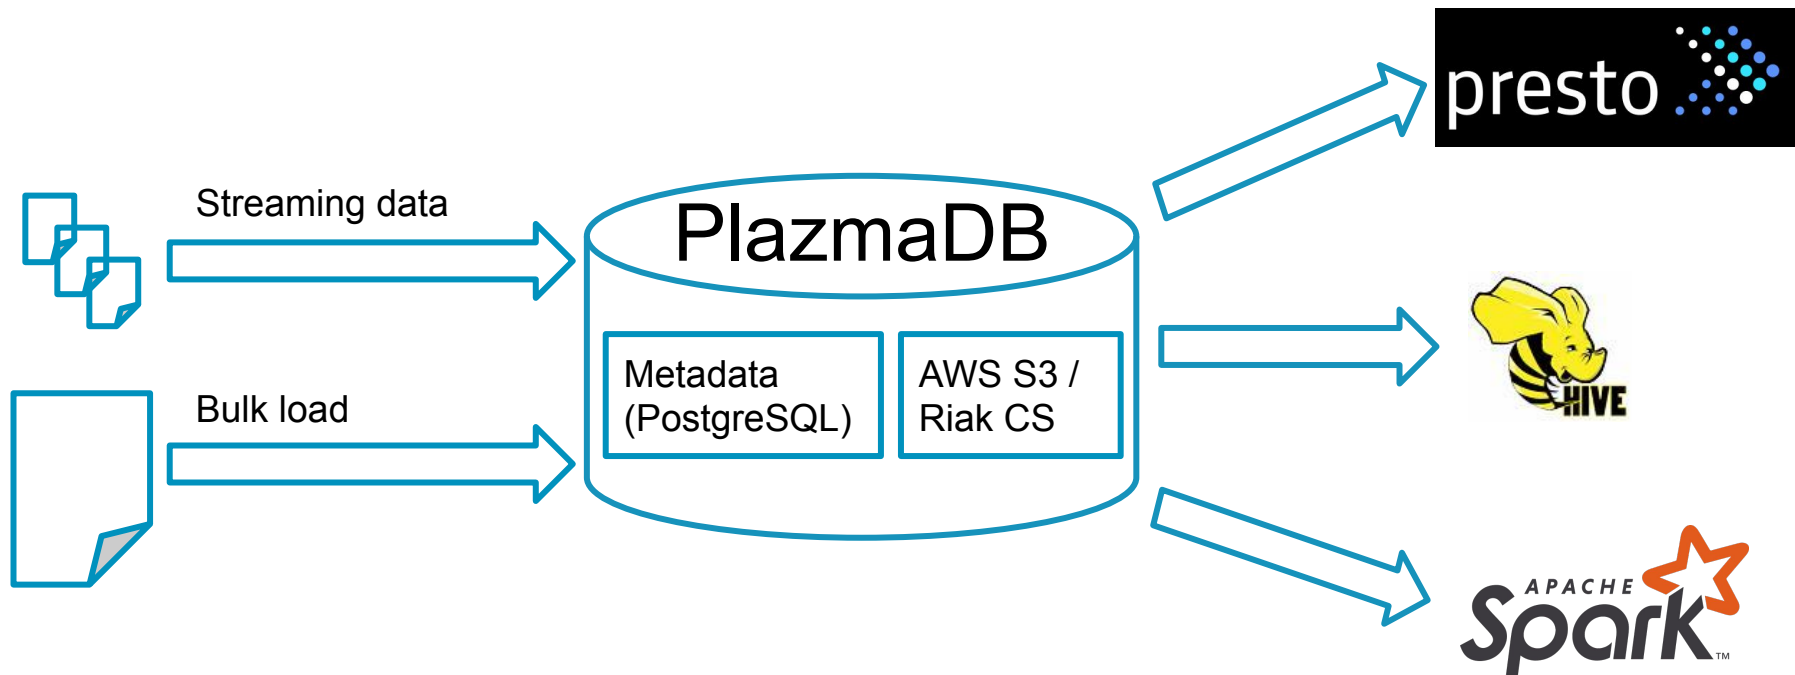


Treasure Data & PlazmaDB

Arm Treasure Data eCDP



PlazmaDB



Daily Workload & Storage Size

Import

500 Billion Records / day
~ 5.8 Million Records / sec

Query

600,000 Queries / day
15 Trillion Records / day

Storage size

5 PB (+5~10 TB / day)
55 Trillion Records

Data Volume

PlazmaDB

Meta DB (PostgreSQL)

Realtime Storage

GiST

Partition
Metadata


Archive Storage

GiST

Partition
Metadata

1 TB

AWS S3 / Riak CS



5 PB

Scaling Strategy

PlazmaDB

Meta DB (PostgreSQL)

Realtime Storage

Archive Storage

Manual Scale up if available
If not, Manual Scale out (sharding)?

Metadata

Metadata

1 TB



AWS S3 / Riak CS



Automatic Scale out

5 PB


Current MetaDB CPU / IO utilization




Q: What is MetaDB capacity?

A: Nobody knew

If PlazmaDB is down..




If PlazmaDB down..

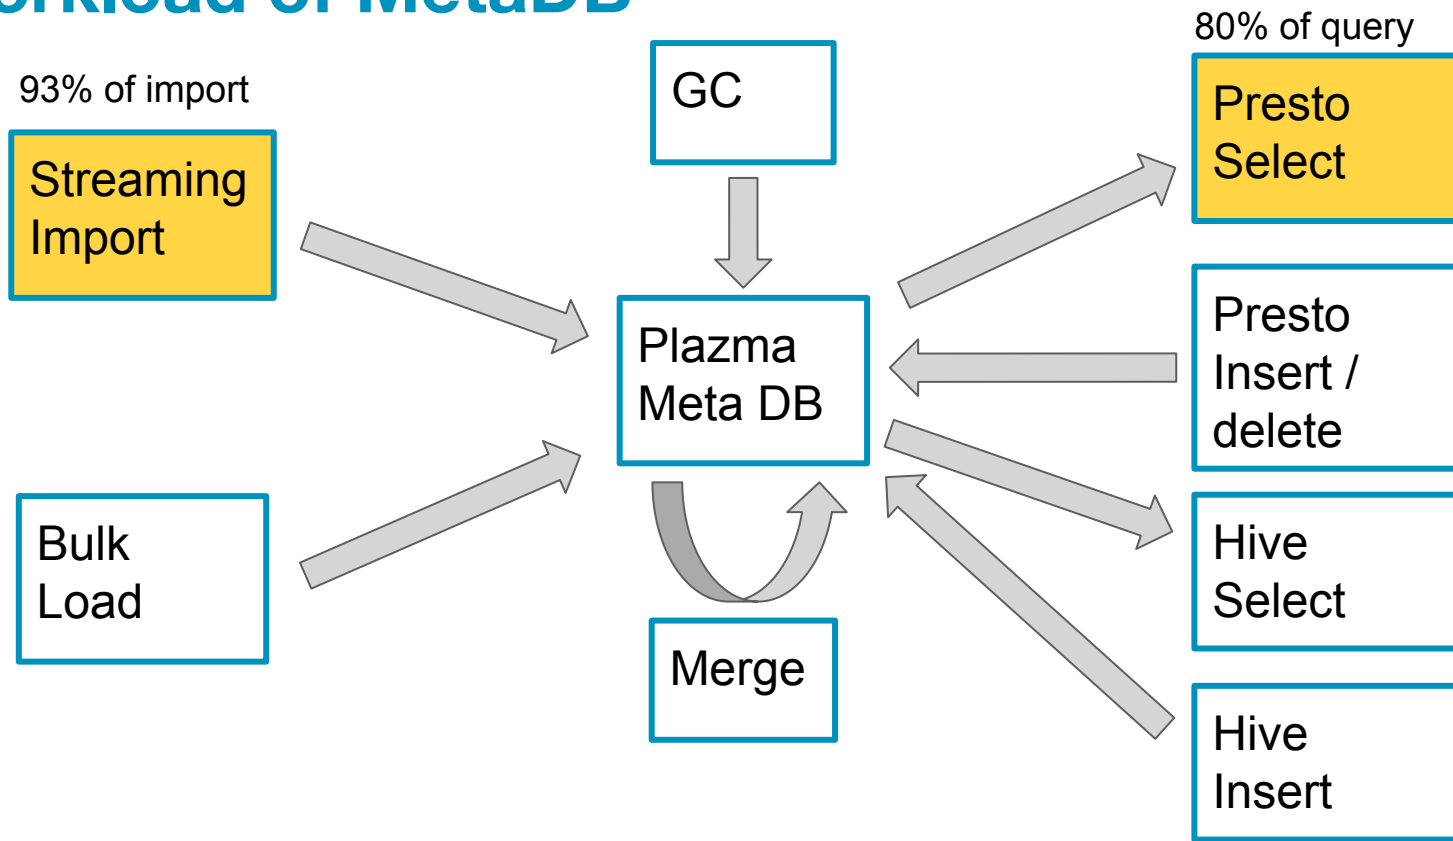


Benchmarking Plazma MetaDB

Workload of MetaDB



Workload of MetaDB



PlazmaDB Streaming Import

Table is collection of partitions

Application

{“time”: “2018-01-01 10:00:00”, “orderid”: 1, ...},
{“time”: “2018-01-01 10:03:03”, “orderid”: 2, ...}

{“time”: “2018-01-01 10:23:03”, “orderid”: 3, ...},
{“time”: “2018-01-01 10:23:12”, “orderid”: 4, ...}

{“time”: “2018-01-01 11:04:44”, “orderid”: 5, ...}

Send logs periodically

EP

Worker

Convert to columnar
& Store a partition

PlazmaDB

2018-01-01 10:00:00	1	1	...
2018-01-01 10:03:03	2	7	...

2018-01-01 10:23:03	3	6	...
2018-01-01 10:23:12	4	3	...

2018-01-01 11:04:44	5	1	...
---------------------	---	---	-----

Partition file: A S3/RiakCS Object

PlasmaDB Metadata

PlasmaDB is Multi tenant

data_set_id: ID combination of User, Database, Table

Meta DB (PostgreSQL)

data_set_id	path	...
1		
1		
1		
2		

AWS S3 / RiakCS

Data set 1

2018-01-01 10:00:00	1	1	...
2018-01-01 10:03:03	2	7	...
2018-01-01 10:23:03	3	6	...
2018-01-01 10:23:12	4	3	...
2018-01-01 11:04:44	5	1	...

Data set 2

2018-01-01 10:00:00	1	1	...
2018-01-01 10:03:03	2	7	...

Partition Index

Meta DB (PostgreSQL)

data_set_id	time_range	path	...
1	[2018-01-01 10:00:00, 2018-01-01 10:03:03]		
1	[2018-01-01 10:23:03, 2018-01-01 10:23:12]		
1	[2018-01-01 11:04:44, 2018-01-01 10:04:44]		
2			

AWS S3 / RiakCS

Data set 1

2018-01-01 10:00:00	1	1	...
2018-01-01 10:03:03	2	7	...
2018-01-01 10:23:03	3	6	...
2018-01-01 10:23:12	4	3	...
2018-01-01 11:04:44	6	1	...

Data set 2

2018-01-01 10:00:00	1	1	...
2018-01-01 10:03:03	2	7	...

Partition Lookup on Analytical Query Processing

Meta DB (PostgreSQL)

data_set_id	time_range	path	...
1	[2018-01-01 10:00:00, 2018-01-01 10:03:03]		
1	[2018-01-01 10:23:03, 2018-01-01 10:23:12]		
1	[2018-01-01 11:04:44, 2018-01-01 10:04:44]		
2			

```
SELECT
  region,
  SUM(price)
FROM
```

```
  orders -- assume this is data set 1
WHERE TD_TIME_RANGE(time,
  '2018-01-01 10:00', '2018-01-01 11:00')
GROUP BY
  region
```

Real World Application is Complicated..

Many Performance Related Factors

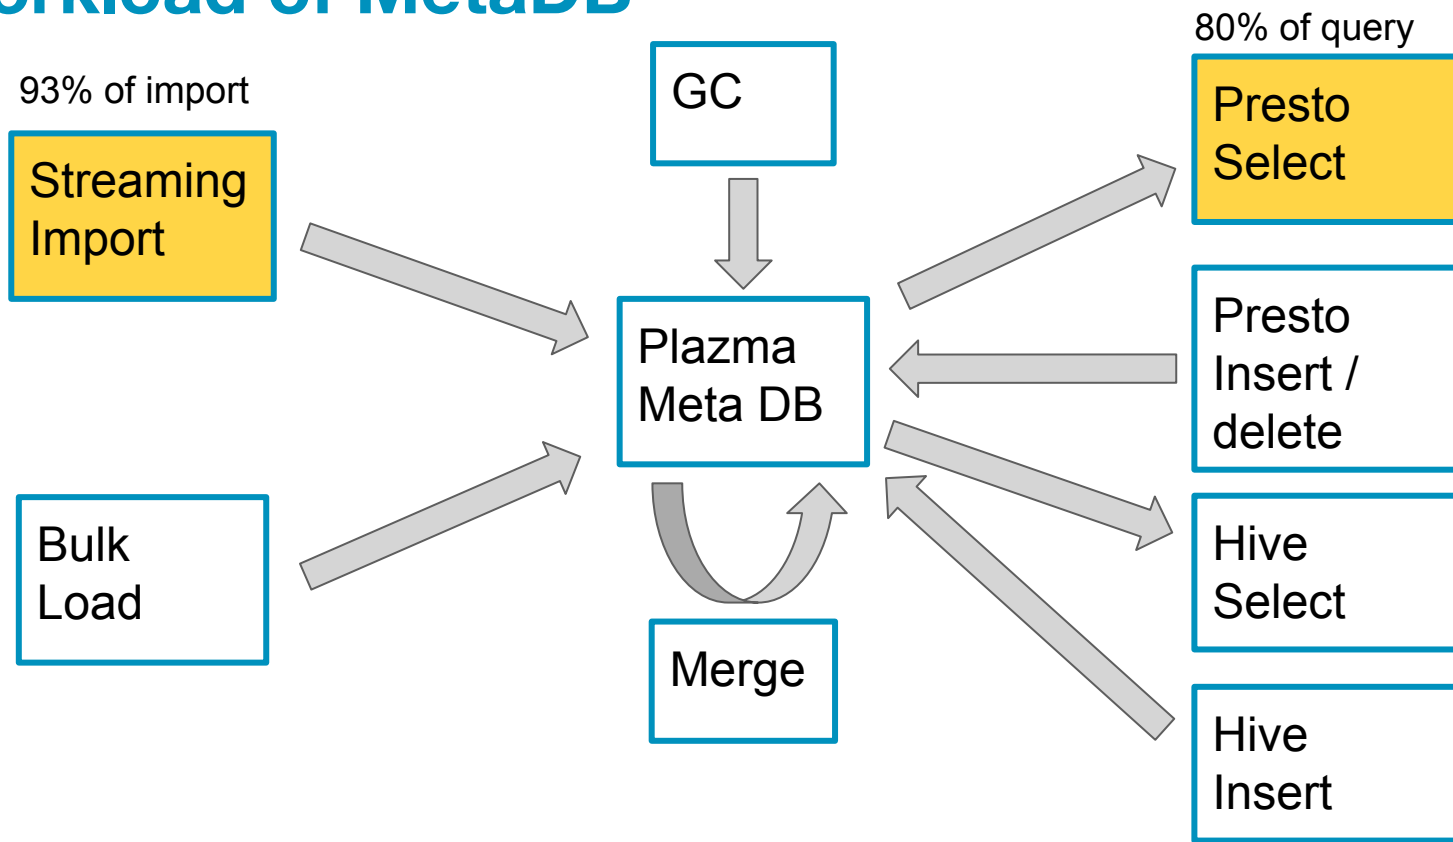
- Type of Workload
- Users' behaviour
 - # of data sets (tables)
 - # of import data
 - # of analytical query request
 - Data skew
- Metadata Storage size
- Server Size (CPU cores, RAM, Storage, Network)
- etc..

Goal of Benchmarking

Define the end of Benchmark task

- Capacity Planning
- Regression Test
- Compare Performance
- Parameter Tuning
- etc..

Workload of MetaDB



Benchmarking Streaming Import

Model of Streaming Import Workload

data_set_id	path	...
3		

Insert a partition metadata



Meta DB (PostgreSQL)

data_set_id	path	...
1		
2		
3		
2		

Performance related factors

Concurrency	(1, 2, 4, 8, 16, 32, 64, 128, 256)
Size of tuple	random based on normal distribution (185 byte on average)

Benchmarking Environment

- AWS RDS PostgreSQL

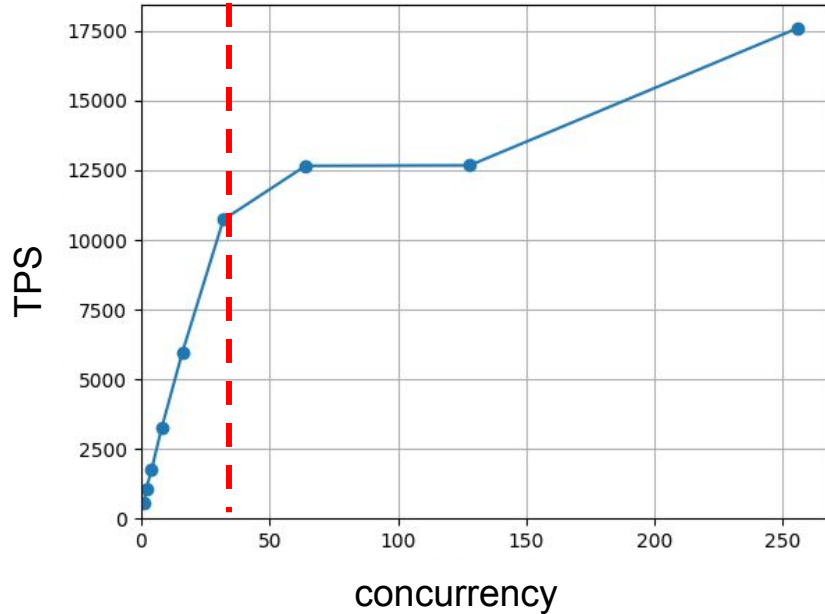
Instance type	db.r3.x8large (32 vcores, 244GB RAM)
Provisioned IOPS	4k
PostgreSQL version	9.4.17

- PostgreSQL parameters

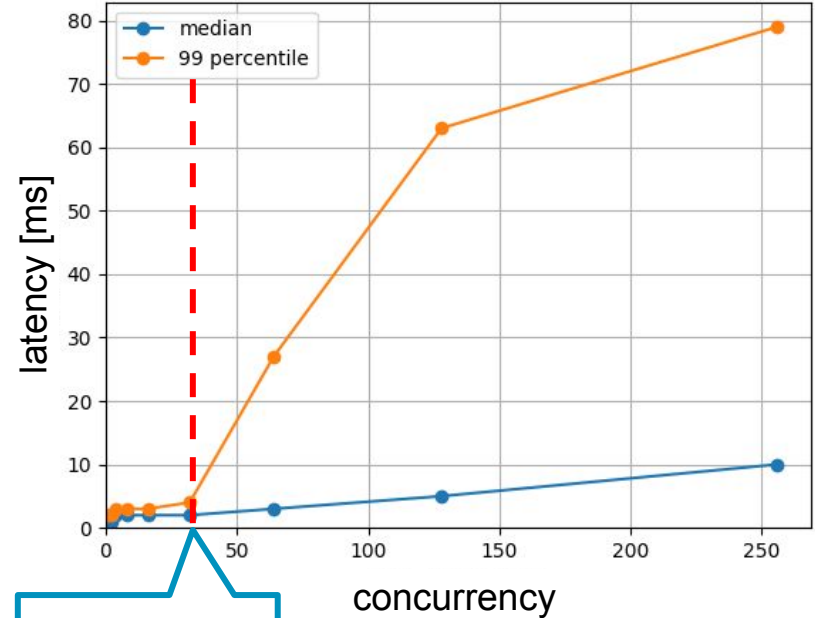
shared_buffers	160GB (~ 60% of RAM)
checkpoint_segments	1500 (24GB)

Scalability of Streaming Import Workload

Throughput



Latency



of cores

Resource Consumption

CPU utilization



Write IO throughput




Write IO


Write IO throughput



Write IOPS



When Write IO issued?




Concurrency and Write IO


Concurrency = 1



Concurrency = 2




Concurrency = 4




IO aggregated

Bottleneck

CPU utilization



Write IOPS



Summary of Streaming Import Workload

- CPU bottleneck
 - Scale almost linearly when concurrency is less than # of cores
 - Throughput can increase after that, but tail latency increases as well
 - Write IOPS doesn't increase as increasing concurrency because of IO aggregation

Benchmarking Presto Select

Model of Presto Select Workload

Meta DB (PostgreSQL)

data_set_id	time_range	path	...
1	[18-01-01 10:00, ... 11:00]	a	
2	[18-01-01 10:00, ... 11:00]	b	
1	[18-01-01 11:00, ... 12:00]	c	
3	[18-01-01 13:00, ... 14:00]	d	
2	[18-01-01 13:30, ... 14:00]	e	
1	[18-01-01 16:00, ... 17:00]	f	
...			
1	[18-01-02 03:00, ... 04:00]	m	

data_set_id=1 and
time_range &&
[18-01-01 00:00, 18-01-02 00:00]



path
a
c
f


Index scan for
data_set_id and
time_range

Performance related factors

Concurrency	(16, 32, 64, 128, 256)
Metadata size	600GB (Dummy data based on actual trend)
# of data sets	30k
Time range to scan (selectivity)	(next slide)
Distribution of data set access frequency	(next slide)


Selectivity

- Random sampling from actual selectivity distribution e.g.)
 - 40% queries: $sl = 1$
 - 5% queries: $sl = [0.01, 0.5]$
 - 5% queries: $sl = [0.5, 0.99]$




Distribution of Data Set Access Frequency

- Metadata size = 600GB
- Shared Buffer size = 160GB
- But, Hot Data size is smaller than Shared Buffer e.g.)
 - 85% of workload comes from 1% data sets
 - 95% of workload comes from 5% data sets




Scalability of Presto Select Workload

Throughput



Latency




Resource Consumption (Concurrency=128)

CPU utilization [%]



Read IO throughput [MB/s]

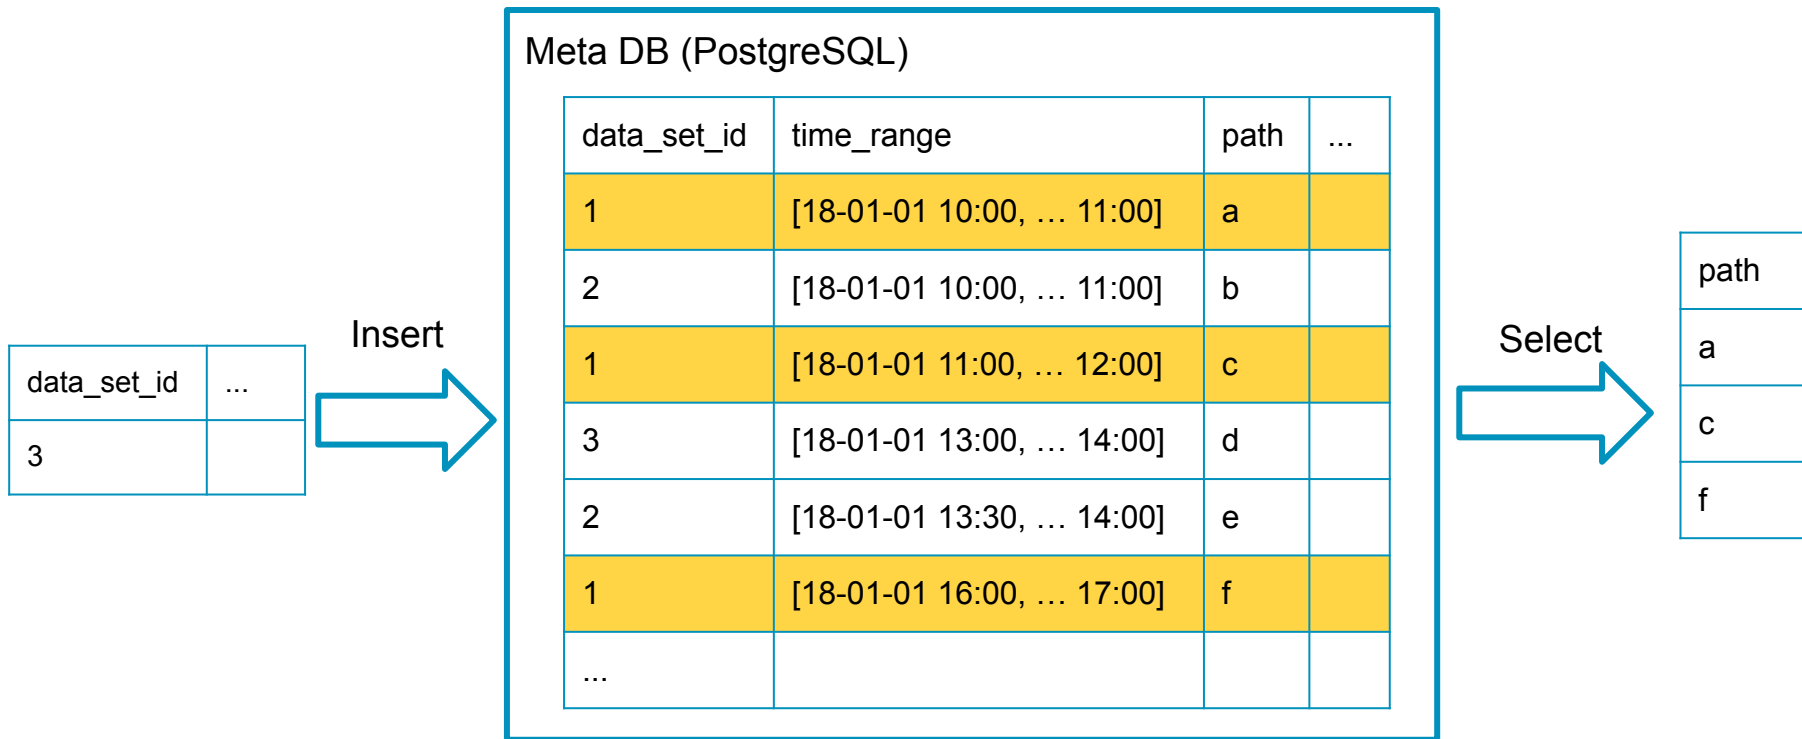


Summary of Presto Select Workload

- CPU bottleneck
 - Scale almost linearly when concurrency is less than # of cores
 - Throughput decreases when concurrency is higher than # of cores
 - Hot data is small enough to fit into DB shared buffer


Benchmarking Mixed Workload

Mixing Streaming Import & Presto Select




Insert : Select = 200 : 1

Realtime Storage & Archive Storage



Realtime Storage & Archive Storage



- Reduced to **1/20** - 1/100 partitions
- Merge can be delayed 5 - **7** hours
- > Metadata will be compressed but accumulate during delay

What is expectation?

- E.g. when concurrency = 64, what throughput will be?
 - Both Streaming Import and Presto Select were CPU bottleneck
- Ref: Single Workload Throughput
 - Streaming Import = 12500 tps
 - Presto Select = 95 tps
- Expectation
 - Streaming Import ~ 6000 tps ?
 - Presto Select ~ 45 tps ?

Result of mixed workload

- Throughput when Concurrency = 64


	Streaming Import [tps]	Presto Select [tps]
mix concurrency=64	5547	27.6
Ref: single	12500	95

Resource Consumption

CPU utilization [%]



Read IO throughput [MB/s]



Bottleneck is changed to Disk IO

Increase PIOPS

	PIOPS	Streaming Import [tps]	Presto Select [tps]
mix concurrency=64	4k	5488	27.4
mix concurrency=64	20k	7179	35.9
Ref: single	4k	12500	95

Cache (DB Shared Buffer) Hit Ratio

	PIOPS	Streaming Import [tps]	Presto Select [tps]	RT storage size	Cache Hit Ratio
mix con=64	4k	5488	27.4	57GB	93%
mix con=64	20k	7179	35.9	75GB	89.5%
Ref: single	4k	12500	95		

DB: shared buffer pages

Partition import 1k/sec

Partition import 2k/sec

Archive Storage

Realtime Storage

Occupy twice shared buffer space

Legend:

- 0.00026M
- 0.0005M
- 0.0039M
- 1.07M
- 0.0021M
- 15.83M

Avg: 0.00026M

Avg: 0.0005M

Avg: 0.0039M

Avg: 1.69M

Avg: 0.0025M

Avg: 16.87M

Impact of Cache Miss

- Avg # of selected rows per a Presto Select = 8000
- A postgres page mostly includes only 1 row for a data set
 - A page (8kB) has 20 - 30 rows, but different data sets' data are stored together
- # of pages to scan per a Presto Select = $8000 / 1 = 8000$
- # of pages to scan per second = $8000 * \text{TPS}$
 - E.g. $\text{TPS}=35 \rightarrow 8000 * 35 = 280\text{k pages/sec}$
 - $\rightarrow 1\% \text{ cache miss causes } 2800 \text{ IOPS}$

Another factor: Auto Vacuum


- Vacuum read cold data
 - Especially Vacuum Freeze does force full scan (before postgres 9.6)
- IO related Vacuum parameters

autovacuum_vacuum_cost_limit	400
autovacuum_vacuum_cost_delay	10ms
vacuum_cost_page_miss	10
vacuum_cost_page_dirty	20

- Max read IOPS = (max IO per vacuum) * (vacuum invoked per sec)
= (400/10) * (1000/10) = 4000

Vacuum is tax of PostgreSQL

Someday auto vacuum stopped accidentally ..



To reduce IO

- Improve cache hit ratio
 - More RAM -> scale up / sharding
 - Improve locality
 - > Partitioning by data set ID: include more relevant rows in a (postgres) page
- Vacuum
 - Avoid full scan of vacuum freeze by using postgres 9.6 or newer

Summary

- IO seems to be a bottleneck
 - Cache miss increases IO dramatically
 - When throughput was increased, we need to pay more tax (vacuum)
- Now we understand what is likely to be a bottleneck
 - Better prioritization of action items
 - Predictable PlazmaDB performance

Don't need to worry about spike and increasing demand :)

Thank You!

Danke!

Merci!

谢谢!

Gracias!

Kiitos!

